

Poor Parallel Program Performance?

Bothersome Bad Bugs?

Sloppy, Slow, Software?

Robert J. Harrison

Stony Brook University

When writing any program


- Correctness is the primary concern
 - Parallel programs have creative ways of going wrong
- Why are you writing a parallel program?
 - To use the resources (memory, compute power, ...) of the parallel computer to solve a big problem fast
- For parallel programs, performance is a correctness issue

Getting it right the first time – Design

- What are your objectives?
 - What sort of parallel computer?
 - How big a problem?
 - How fast a solution?
 - How many processors?
 - How portable?
- You can reason about program performance by constructing a performance model


If you have an infinitely fast processor ...

- How fast can you compute?
 - As fast as you can get data to it
- Data motion is more often the bottleneck than how fast computation actually happens
 - FLOPs are free
 - Bandwidth is expensive (power and price)



Latency and bandwidth

- Latency – startup time
 - Cost of operating on zero elements
- Bandwidth – steady state #elements / unit time



$$t(n) = t_0 + n/B$$

$$rate = \frac{n}{t} = \frac{B}{B t_0 / n + 1} \Rightarrow n_{1/2} = B t_0$$

Matrix multiplication on a simple sequential computer - I

```
for i=1 to N
```

```
  for j=1 to N
```

```
    sum = 0.0
```

```
    for k=1 to N
```

```
      sum += A[i, k] * B[k, j]
```

```
    C[i, j] = sum
```

$$c_{ij} = \sum_{k=1}^N a_{ik} b_{kj}$$


$2N^3$ FLOPs

$2N^3$ memory loads

Compute
8 GFLOP/s
Memory bandwidth
1 GWord/s
Memory latency
0.2 to 2us

A read with unit stride
B read with stride N

Disaster!



Matrix multiplication on a simple sequential computer - II

```
for i=1 to N
```

```
  for k=1 to N
```

```
    aik = A[i,k]
```

```
    for j=1 to N
```

```
      C[i,j] += aik * B[k,j]
```

$$c_{ij} = \sum_{k=1}^N a_{ik} b_{kj}$$


$2 N^3$ FLOPs

$3 N^3$ memory load + store

Compute
8 GFLOP/s
Memory bandwidth
1 GWord/s
Memory latency
0.2 to 2us

B read with unit stride
C written with unit stride

12.5% of peak speed



Matrix multiplication on a simple sequential computer - III

```
for i=1 to N
```


```
  for k=1 to N
```

```
    aik = A[i,k]
```

```
    for j=1 to N
```

```
      C[i,j] += aik * B[k,j]
```

$$c_{ij} = \sum_{k=1}^N a_{ik} b_{kj}$$



Compute

8 GFLOP/s

Cache bandwidth

8 GWord/s

Cache latency

0.5ns

Memory bandwidth

1 GWord/s

Memory latency

0.2 to 2 μ s

$2N^3$ FLOPs

$3N^3$ memory load + store

B read with unit stride

C written with unit stride

66% speed ... iff B&C small enough to fit in cache

Matrix multiplication on a simple sequential computer - IV

```
for tiles of i
```

```
  for tiles of j
```


```
    for k=1 to N
```

```
      for i in tile       $c_{ij} = \sum_{k=1}^N a_{ik} b_{kj}$ 
```

```
        aik = A[i,k]
```

```
        for j in tile
```

```
          C[i,j] += aik*B[k,j]
```



Choose tile size T so that $2 \cdot T \cdot T$ fits into cache
Usually must do this empirically

Memory motion and compute take the same time

Assuming compute and memory motion can be overlapped we are now running at 66% peak speed

Robert J. Harrison

To get peak speed must keep C in registers

Matrix multiplication on a simple sequential computer - V

```
for tiles of i
```

```
  for tiles of j
```

```
    for k=1 to N
```

```
      for i in tile
```

```
        aik = A[i,k]
```

```
        load C[i,*] into R[*]
```

```
        for j in tile
```

```
          R[j] += aik*B[k,j]
```

```
        store R[*] into C[i,*]
```

P



Cache



Memory

Now we are running at full speed

Depressing conclusion is that we must tile for every level of the memory hierarchy

Parallel Algorithms

- All of the complexity of sequential algorithms
 - Plus ...
- Concurrency
 - Gotta keep everyone equally busy
- Additional levels in the memory hierarchy
 - Cache and memory of other processors
- Contention
 - Literally processes fighting over wires

Weak and strong scaling

- Strong scaling
 - Fixed problem size as increase number of processors
- Weak scaling
 - Increase the problem size as increase number of processes

Speedup, efficiency, iso-efficiency

$$S(P) = \frac{t(1)}{t(P)} \quad S_{\text{ideal}}(P) = P$$

$$\epsilon(P) = \frac{S(P)}{P} = \frac{t(1)}{P t(P)} \leq 1 \quad \epsilon_{\text{ideal}}(P) = 1$$

$t(P, N, \epsilon) =$ isoefficiency : running time of problem size N required to get specified efficiency on P processors
... for perfect weak scaling is a constant
... for perfect strong scaling is a decreasing function of P

Amdahl's Law

$$t(P) = t_{seq} + \frac{t_{par}}{P}$$


- Assume
 - Sequential component (or equivalently a part that each process must repeat)
 - Perfectly parallel component
- Maximum speed up is

$$S(P) = \frac{t(0)}{t(P)} = \frac{t_{seq} + t_{par}}{t_{seq} + \frac{t_{par}}{P}}$$

$$S(\infty) \rightarrow \frac{t_{par}}{t_{seq}}$$

50% efficiency when $P = \frac{t_{par}}{t_{seq}}$

Amdahl's law in action



$$t(p) = 10 + \frac{1000}{P}$$

Contention

- Multiple processes trying to
 - Read the same memory location or bank
 - Send a message over the same wire
 - Compute on the same processor
- Execution is serialized
 - Equivalent to reducing execution rate
 - E.g., P processes trying simultaneously to send data over a link of bandwidth B
 - Each process sees bandwidth B/P

New failure modes

- Race conditions
 - Read before/after write
 - Concept of atomicity
- Dead/live lock
- Hangs
- Exhausting buffer space
- Oversubscribed resources
 - Exponential slowdown

Safe programming models

- Some models are designed to eliminate some types of errors
 - Communicating sequential processes (CSP, Hoare)
 - MPI safe message passing
 - Bulk synchronous program (BSP)
 - OpenMP data parallel