

SICM² Software Summer School: Parallel algorithms and programming models

Robert Harrison¹

¹ Stony Brook University

July 2014
Stony Brook, NY

High-performance computing — objectives

- Do science using the computer
 - Emphasizes productivity, speed of testing new ideas, correctness
- Science is collaborative, funding limited, software long lived
 - Need readable, extensible, maintainable, portable software
- Science is done by science students not computer scientists
 - Ideally code in high-level concepts, hiding arcane computer details
- Big/many calculations and/or a large user community
 - Demand high-performance, efficient use of resources, robustness

Practical high-performance computing

- Performance is a level-1 correctness issue
- How to make a computer run fast?
- How to do this portably?
 - Current machines
 - Future machines
- How to do this while maintaining productivity?

Technology trends

- Power wall
 - Computer clock frequencies are no longer steadily increasing
- Memory wall
 - Memory, communication, and disk bandwidth increasing less rapidly than compute speed
 - New technologies (e.g., stacked memory) essential to progress
- Parallelism is the **only** path to increased performance
 - Fine-grain
 - instructions — MIMD
 - vectors — SIMD
 - Medium-grain
 - cores — MIMD
 - Coarse-grain
 - sockets, nodes — MIMD
- Good news — technology is delivering more fine/medium grain
 - Phew! 1M-way coarse grain parallelism is too much!
 - What to do with all those transistors?

Connecting concepts in architecture, software, and tools

Architectural element	Software construct	Programming tool
<i>within a "core"</i> pipe-lined units multiple instruction simd units	complex expressions, loops complex expressions, fat loops loops, ops on vectors/matrices	compiler, library compiler compiler, library
<i>within a "node"</i> cores/thread units accelerators	thread/process/task loop nests, ops on matrices task, ops on matrices	openmp, pthreads, tbb openmp, opencl, libraries opencl, openacc, cuda, libs
<i>between "nodes"</i> distributed memory	outer loops, multiple tasks, big matrices	mpi, libraries

Just do it!

- Overthinking all this is counter productive
- Write pipelineable, vectorizable code — more on this later
- Have the compiler do the work

```
for (int i=0; i<n; i++) {  
    c[i] = a[i]*23.0 + exp(d[i]*d[i]);  
}
```

Hello world vector style

■ Code

```
#include <iostream>

int main() {
    double a[100];
    for (int i=0; i<100; i++) a[i] = i;
    double sum = 0.0;
    for (int i=0; i<100; i++) sum += a[i];
    std::cout << sum << std::endl;
    return 0;
}
```

■ Compilation

```
$ icpc -fast -vec_report3 sum.cc
sum.cc(5): (col. 5) remark: FUSED LOOP WAS VECTORIZED.
```

■ Output ($\sum_{i=0}^{n-1} i = n(n-1)/2 = 4950$)

```
$ ./a.out
4950
```

Thinking about performance

- What is the bottleneck
 - Data motion?
 - Instruction issue?
 - Functional units?
- Moving data
 - disk
 - main memory
 - some level of cache
 - registers
- Performance model

$$t(n) = L + \frac{n}{B}$$

L = latency, B = bandwidth, n = number of elements

Hockney's $n_{1/2}$

- How long must a loop be, or how much data must we move, in order to hit 50% performance

$$\text{rate}(n) = \frac{n}{t(n)} = \frac{n}{L + \frac{n}{B}} = \frac{B}{\frac{BL}{n} + 1}$$

- The asymptotic rate is just

$$\text{rate}(\infty) = B$$

- Hence,

$$\text{rate}(n_{1/2}) = \frac{B}{2} \rightarrow n_{1/2} = BL$$

- also

$$\text{rate}(n_{90\%}) = 0.9B \rightarrow n_{90\%} = 9BL$$

E.g., typical message passing between processes

- latency 5us, bandwidth 1e9 bytes/s
- $n_{1/2} = 1,000$ bytes
- $n_{90\%} = 9,000$ bytes

Pipelined functional units

- Parallelism arises from overlapping stages of performing successive operations
- E.g., floating point $a*b$
 - A single operation may take 3 cycles to complete (pipeline depth or latency)
 - You can issue a new operation every cycle

cycle	stage0	stage1	stage2	result
1	$a0*b0$			
2	$a1*b1$	$a0*b0$		
3	$a2*b2$	$a1*b1$	$a0*b0$	
4	$a3*b3$	$a2*b2$	$a1*b1$	$a0*b0$
5	$a4*b4$	$a3*b3$	$a2*b2$	$a1*b1$
⋮	...			

Pipelined functional units

- Serial execution

- Wait for each result to complete
- The cost per operation is the pipeline depth

- Pipelined execution

- $t(n) = L + n - 1$ (in cycles)
- $n_{1/2} = L - 1$
- $n_{90\%} = 9 * (L - 1)$ $9 * 2 = 18$

Multiple instruction issue

- High-end x86 can issue each cycle at least
 - One or more integer operations (e.g. inc loop counter)
 - A test+branch operation
 - A (simd) floating point multiply
 - A (simd) floating point addition
 - Two (simd) memory operations
- The integer units typically have 1 cycle latency
- FMA — fused multiply-add — $a*b+c \rightarrow d$
 - Introduced with Haswell (<http://goo.gl/Bo3Rt>, <http://goo.gl/tePyD>)
- The compiler does this for you — usually better than humans except on very tight loops.

SIMD (vector) operations

- Single instruction applying operation to multiple data

$$\begin{pmatrix} a_0 \\ a_1 \\ a_2 \\ a_3 \end{pmatrix} * \begin{pmatrix} b_0 \\ b_1 \\ b_2 \\ b_3 \end{pmatrix} \rightarrow \begin{pmatrix} a_0 * b_0 \\ a_1 * b_1 \\ a_2 * b_2 \\ a_3 * b_3 \end{pmatrix}$$

- Older x86 — SSE — 2 doubles (128 bit)
- Newest x86 — AVX(2) — 4 doubles (256 bit), predicates, special functions, gather, scatter
- MIC x86 — LRBNI — 8 doubles (512bit), predicates, special functions, gather, scatter
- Zillions of different operations (<http://goo.gl/D0d7t>)
- Operations are pipelined
- $n_{90\%} = 9 * W * (L - 1) = 108$ with $W=4$ (width) and $L=4$ (latency)

Serial scalar v.s. pipelined SIMD performance

- Serial scalar computation produces 1 result every L cycles
- Pipelined SIMD produces W results every cycle
- Pipelined is WL times faster which is 16 with $W = L = 4$
- Allowing for 2 operations per cycle ($*, +$) this becomes 32
- I.e., worst case serial, non-vector code can be 32x slower!!!!
- Whenever some twit talks about GPGPUs being 100+ times faster than a CPU ask about how well vectorized (and threaded) the CPU code was.

Example performance — DAXPY

- Measure cycles/iteration of this loop as a function of n

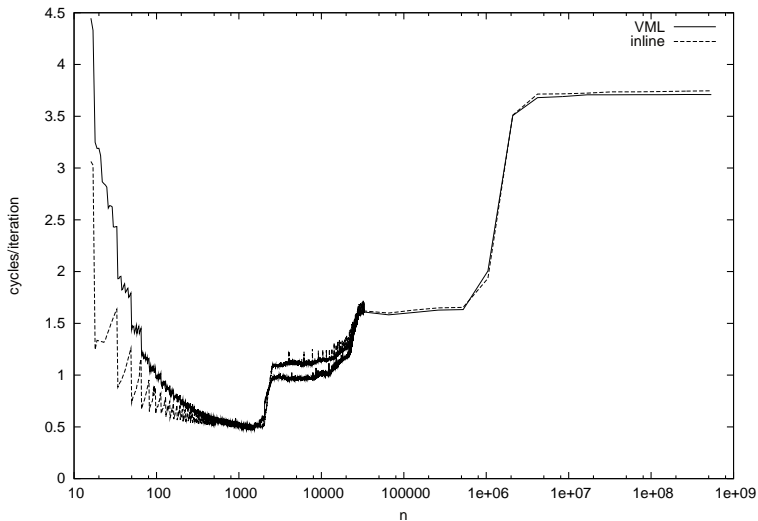
```
for (int i=0; i<n; i++) {  
    y[i] = a*x[i] + y[i];  
}
```

or from VML

```
cblas_daxpy (n, a, x, 1, y, 1);
```

- What are you expecting to see?

Example performance — DAXPY



DAXPY performance analysis

- In two cycles can do 2 loads, *, +, store, increment, branch
- AVX SIMD is 4 wide
- Hence, 4 iterations per 2 cycles = 0.5 cycles/iter at best
- L1 = 32KB (4096 dbl), 48 bytes/cycle
(capable of 64 bytes/cycle?)
- L2 = 256KB (32768 dbl), 21 bytes/cycle
- L3 = 8MB (1M dbl), 14.8 bytes/cycle
- main memory, 6.4 bytes/cycle

DAXPY performance analysis

- Lesson — cache fast, memory slow (latency even more so)

- Assignment — predict then measure performance of

```
for (int i=0; i<n; i+=4) {  
    y[i] = a*x[i] + y[i];  
}
```

(remember to account for reduced number of elements in perf. anal.)

- Assignment — what if stride is 4096 instead of 4?

Read the documentation — really!

- The tools do a **lot** more than most people realize.
- Will introduce lots of valuable concepts and techniques.
- Read the release notes with every new version.
- Intel developer zone <http://software.intel.com/>
- Intel compilers <http://software.intel.com/en-us/intel-compilers>
- Intel composer software suite
<http://software.intel.com/en-us/articles/intel-c-composer-xe-2011-documentation>
- Intel C++ compiler <http://software.intel.com/sites/products/documentation/doclib/stdxe/2013/composerxe/compiler/cpp-lin/index.htm>
- Getting started with auto-vectorization http://software.intel.com/sites/products/documentation/doclib/stdxe/2013/composerxe/tutorials/lin/cmp_vec_c/index.htm

More Intel links

- Go parallel portal <http://go-parallel.com/>
- Learning lab <http://software.intel.com/en-us/intel-learning-lab/>
- Vectorization <http://software.intel.com/en-us/intel-vectorization-tools/>
- Evaluation guide <http://software.intel.com/en-us/evaluation-guides/>
- Parallel magazine <http://software.intel.com/en-us/intel-parallel-universe-magazine/>
- OpenCL <http://software.intel.com/en-us/vcsource/tools/opencl-sdk>
- MxM using MKL http://software.intel.com/sites/products/documentation/doclib/mkl_sa/11/tutorials/mkl_mmx_c/tutorial_mkl_mmx_c.pdf

Non-vectorizable loops

- Dependencies between iterations

```
for (int=1; i<n; i++) a[i-1] = 3*a[i];
```

- Dependencies through memory

```
void f(int n double *a, double* b) {  
    for (int=0; i<n; i++) a[i] = b[i];  
}
```

- Dependencies through iteration index or count

```
for (int i=0; i<n; i++) {if (a[i] > 1) n++;}  
for (int i=0; i<n; i++) {if (a[i] > 1) i++;}
```

- Calls to non-inline functions (math library exceptions)
- If-tests that don't resolve to vector merge

Passing your knowledge to the compiler

- `#pragma ivdep` — ignore vector dependencies
- `restrict`
- `-fargument-noalias` — function arguments cannot alias each other, but they can alias global storage
- `-fargument-noalias-global` — function arguments cannot alias each other or global storage
- `#pragma loop count (n)` — advises the compiler of the typical trip count of the loop.
- `#pragma vector always` — always vectorize if safe regardless of if performance improvement expected
- `#pragma vector align` — asserts that data within the following loop is aligned (to a 16 byte boundary, for SSE instruction sets)
- `#pragma novector` — asks the compiler not to vectorize a particular loop.

Passing your knowledge to the compiler

- `#pragma vector nontemporal` — gives a hint to the compiler that data will not be reused, and therefore to use streaming stores that bypass cache.
- `#pragma simd` — Lots of options. Forces vectorization even if it is unsafe.
- `__attribute__((vector))` and `__declspec(vector)` — Vectorization of functions when definition is not available for inlining.

Monte Carlo Example

- Metropolis Monte Carlo
 - General and powerful algorithm for multi-dimensional integration
 - Abuse it to create a simple test code that reflects real applications
- Compute average value of x sampled from the normalized probability distribution function e^{-x} , i.e.,

$$\langle x \rangle = \frac{\int_0^{\infty} x e^{-x} dx}{\int_0^{\infty} e^{-x} dx} = 1$$

Monte Carlo Example - II

- Starting from uniform random numbers in $[0, 1)$ use Metropolis algorithm to sample from e^{-x}
- Approximate infinity as 23 ($\exp(-23) = 1e-10$)
- Algorithm

```
x = 23.0*rand() // initialize
while (1) {
    xnew = 32.0*rand();
    if (exp(-xnew) > exp(-x)*rand()) x = xnew;
}
```

- Asymptotically, x is sampled from $\exp(-x)$ (with some correlation between successive values)

Monte Carlo Example - II

- Intel(R) Xeon(R) CPU E5-2687W
 - mc0 — 71.2 cycles/point
 - mc1 — 71.8
 - mc2 — 76.2
 - mc3 — 54.3
 - mc4 — 18.0
 - mc5 — 14.8
- Intel(R) Xeon(R) CPU E5645 (Hokiespeed)

Variational Monte-Carlo Example

- Uses Metropolis algorithm to sample points distributed according to ψ^2 for the helium atom and the Hylleraas wave function

$$\psi(r_1, r_2, r_{12}) = (1 + \frac{1}{2}r_{12})e^{-2(r_1+r_2)}$$

- Variational since $\langle E_L \rangle \geq E_0$ where $E_L = (\hat{H}\psi)/\psi$
- Computes $\langle r_1 \rangle$, $\langle r_2 \rangle$, $\langle r_{12} \rangle$

Assignments for vector programming

- 1** Read (at least skim) all the compiler reference manual
Pay special attention to the autovectorization section
- 2** Work through the autovectorization getting started
- 3** Examine the model matrix multiplication code
- 4** Understand each of the Monte Carlo code versions and reproduce performance
- 5** Tune a kernel relevant to you (modify bench.cc to measure performance)

Shared-memory programming beyond data-parallel OpenMP

■ OpenMP

<https://computing.llnl.gov/tutorials/openMP/>

<https://iwomp.zih.tu-dresden.de/downloads/omp30-tasks.pdf>

<http://www.cs.utah.edu/~mhall/cs4961f11/>

<http://www.cs.utah.edu/~mhall/cs4961f11/CS4961-L9.pdf>

■ Pthreads

<https://computing.llnl.gov/tutorials/pthreads/>

<http://www.yolinux.com/TUTORIALS/LinuxTutorialPosixThreads.html>

<http://cs.gmu.edu/~white/CS571/pthreadTutorial.pdf>

■ TBB

<http://www.inf.ed.ac.uk/teaching/courses/ppls/TBBtutorial.pdf>

<http://stackoverflow.com/questions/13446434/tbb-beginner-tutorial>

[http://software.intel.com/sites/products/documentation/doclib/tbb_sa/help/index.htm#](http://software.intel.com/sites/products/documentation/doclib/tbb_sa/help/index.htm#reference/reference.htm)

<reference/reference.htm>

Spinlocks v.s. mutexes

■ Spinlocks

- Fastest with low contention and no false sharing
- Low resource use (minimally one cache line)
- Potentially very slow under contention, in virtual machines, oversubscribed
- Spinning threads can saturate memory and slow everything

■ Mutexes

- Thread scheduler blocks waiting threads
- Usually slower than spinlocks in low contention
- Scale well under contention, in virtual machines, oversubscription

■ Recommendations

- Use trusted implementation of higher-level concept (e.g., thread-safe containers)
- Think tasks, not threads
- Avoid heavily contested critical sections
- Use mutexes; try spinlocks if performance is poor

Critical section issues

- Fairness — not guaranteed but often assumed
- Correctly updating shared data structures is not easy
 - `volatile` — does not solve the underlying problem but is useful to get the compiler to help

http://www.hpl.hp.com/personal/Hans_Boehm/c++mm/faq.html

<http://software.intel.com/en-us/blogs/2007/11/30/>

[volatile-almost-useless-for-multi-threaded-programming](#)

<http://www.drdobbs.com/cpp/volatile-the-multithreaded-programmers-b/184403766>

- Need memory fences and compile/run time barriers to instruction migration
- Pthread routines provide these — home grown assembly probably does not
- STL structures cannot be made `volatile` — but may not need to be

Other options

- Condition variables
- Fair, scalable synchronization
<http://dl.acm.org/citation.cfm?id=103729>
- Again, think tasks not threads

Intel Thread Building Blocks